


"Lending Club Analysis and Data Cleaning using PySpark"

Description:

1. In this project I have used the lending club dataset. Lending club is a peer to peer lending company. They will be providing the loans to companies, peoples. Sometimes peoples will not have good credit history but they want immediately want loan will come to lending club and ask for loan.
2. There will be set of peoples who will be seeing the people's profile
 - a. If the applicant is likely to repay the loan, then not approving the loan results in a loss of business to the company
 - b. If the applicant is not likely to repay the loan, i.e. he/she is likely to default, then approving the loan may lead to a financial loss for the company
3. In this the interest rate are higher than the other banks.
4. Above is the overview of the lending club and I have used the the 2018 Q4 dataset.

Dataset:

The size of the dataset is 70mb which contains 128K records & 125 columns.

**Problem:**

The values present in the dataset is uncleaned(Some columns like term contains the datatypes as `string('36 months')` but in the business use we need to convert it into `int(36)`).

If we downstream this data to machine learning models we have to make it to numeric.

Data Cleaning overview:

1. Since it contains 145 columns need to limit it into another df.

2.

```
df_sel.describe("loan_amnt","emp_length","dti","delinq_2yrs","revol_util","total_acc").show()
```

summary	loan_amnt	emp_length	dti	delinq_2yrs	revol_util	total_acc
count	128412	128412	128175	128412	128256	128412
mean	15971.32102139987	null	19.933177530719778	0.22783696227766875	null	22.677413325857398
stddev	10150.384232741915	null	20.143542243475476	0.7337929617806072	null	12.129215673024733
min	1000	1 year	0.0	0	0%	2
max	40000	n/a	999.0	24	99.90%	160

From the above image we can see in the 'emp_length' for the min it is showing as '1 year', mean & stddev it has treated as string but there is an integer value. Also we are having a good outliers for 'dti' for mean & stddev. We need to clean the above highlighted so that we can get the proper statistics.

3.

```
#Checking employee length variable
spark.sql("select distinct emp_length from loanstats limit 50").show()
```

emp_length
5 years
9 years
1 year
n/a
2 years
7 years
8 years
4 years
6 years
3 years
10+ years
< 1 year

From the 'emp_length' we could see 4 different types of values present in it, we need only the values in **numeric** type.

4. Clean the 'term' column
5. Clean the 'revol_util' column.
6. Clean the 'dti' column -> Null values
7. Clean the 'loan_status'

Data Aggregation:

1.Schema of the dataset.

```
df.printSchema()

root
|-- id: string (nullable = true)
|-- member_id: string (nullable = true)
|-- loan_amnt: integer (nullable = true)
|-- funded_amnt: integer (nullable = true)
|-- funded_amnt_inv: double (nullable = true)
|-- term: string (nullable = true)
|-- int_rate: string (nullable = true)
|-- installment: double (nullable = true)
|-- grade: string (nullable = true)
|-- sub_grade: string (nullable = true)
|-- emp_title: string (nullable = true)
|-- emp_length: string (nullable = true)
|-- home_ownership: string (nullable = true)
|-- annual_inc: double (nullable = true)
|-- verification_status: string (nullable = true)
|-- issue_d: string (nullable = true)
|-- loan_status: string (nullable = true)
|-- pymnt_plan: string (nullable = true)
|-- url: string (nullable = true)
|-- desc: string (nullable = true)
|-- purpose: string (nullable = true)
|-- title: string (nullable = true)
|-- zip_code: string (nullable = true)
|-- addr_state: string (nullable = true)
|-- dti: double (nullable = true)
|-- delinq_2yrs: integer (nullable = true)
|-- earliest_cr_line: string (nullable = true)
|-- inq_last_6mths: integer (nullable = true)
|-- mths_since_last_delinq: integer (nullable = true)
```

2. Displaying the records

df.show()

	id	member_id	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	grade	sub_grade	emp_title	emp_length
	null	null	10000	10000	10000.0	36 months	10.33%	324.23	B	B1	null	< 1 year
	null	null	2500	2500	2500.0	36 months	13.56%	84.92	C	C1	Chef	10+ years
	null	null	12000	12000	12000.0	60 months	13.56%	276.49	C	C1	null	< 1 year
	null	null	15000	15000	14975.0	60 months	14.47%	352.69	C	C2	null	n/
	null	null	16000	16000	16000.0	60 months	17.97%	406.04	D	D1	Instructional Co...	5 years
	null	null	9600	9600	9600.0	36 months	23.40%	373.62	E	E1	driver coordinator	9 years
	null	null	4000	4000	4000.0	36 months	23.40%	155.68	E	E1	Security	3 years
	null	null	3500	3500	3500.0	36 months	20.89%	131.67	D	D4	gas attendant	10+ years
	null	null	9600	9600	9600.0	36 months	12.98%	323.37	B	B5	null	n/
	null	null	8000	8000	8000.0	36 months	23.40%	311.35	E	E1	Manager	10+ years
	null	null	16000	16000	16000.0	60 months	26.31%	481.99	E	E4	Financial Relatio...	< 1 year
	null	null	30000	30000	30000.0	60 months	18.94%	777.23	D	D2	Postmaster	10+ years
	null	null	10000	10000	10000.0	60 months	19.92%	264.5	D	D3	Material Handler	10+ years
	null	null	5000	5000	5000.0	36 months	17.97%	180.69	D	D1	Administrative	6 years
	null	null	23000	23000	23000.0	60 months	20.89%	620.81	D	D4	Operator	5 years
	null	null	32075	32075	32075.0	60 months	11.80%	710.26	B	B4	Nursing Supervisor	10+ years
	null	null	13000	13000	13000.0	36 months	23.40%	505.95	E	E1	Sale Representative	2 years
	null	null	35000	35000	35000.0	60 months	15.02%	833.02	C	C3	Sr Sales Manager	10+ years
	null	null	40000	40000	40000.0	60 months	11.31%	875.9	B	B3	null	< 1 year
	null	null	8000	8000	8000.0	36 months	10.33%	259.38	B	B1	Optician	6 years

only showing top 20 rows

3. How many rows are in the dataset.

#check how many rows are there in the dataset
df.count()

128412

4. Creating a temporary variable to run the sql queries

#creating a temporary table to run the SQL queries.
df.createOrReplaceTempView("loansatats")

5. Since there are more than 145 columns limiting only the selected columns

#selecting only the required columns which will be only usable.
df_sel = df.select(["term", "home_ownership", "grade", "purpose",
"int_rate", "addr_state", "loan_status", "application_type",
"loan_amnt", "emp_length", "annual_inc", "dti", "delinq_2yrs",
"revol_util", "total_acc", "num_tl_90g_dpd_24m", "dti_joint"])

6. When displaying the output we can see the below error

```
df_sel.describe("loan_amnt","emp_length","dti","delinq_2yrs","revol_util","total_acc").show()
```

summary	loan_amnt	emp_length	dti	delinq_2yrs	revol_util	total_acc
count	128412	128412	128175	128412	128256	128412
mean	15971.32102139987	null	19.933177530719778	0.22783696227766875	null	22.677413325857398
stddev	10150.384232741915	null	20.143542243475476	0.7337929617806072	null	12.129215673024733
min	1000	1 year	0.0	0	0%	2
max	40000	n/a	999.0	24	99.90%	160

7. Checking the 'emp_length'

```
#Checking employee length variable
spark.sql("select distinct emp_length from loanstats limit 50").show()
```

emp_length
5 years
9 years
1 year
n/a
2 years
7 years
8 years
4 years
6 years
3 years
10+ years
< 1 year

It contains 4 different types of values present in it but we need only the numeric values.

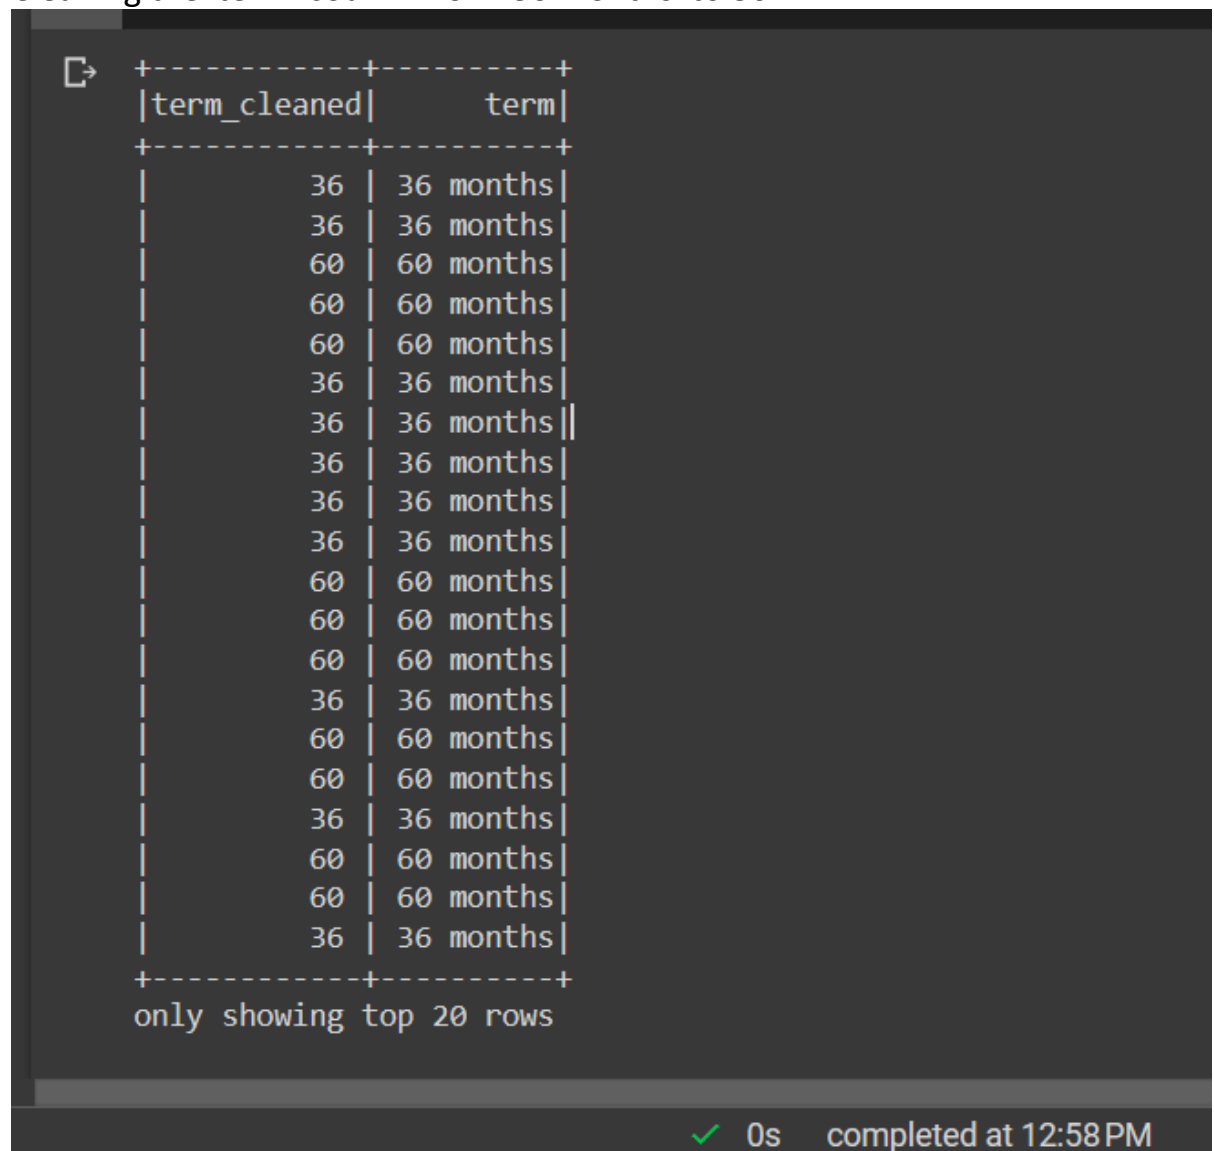
Data after cleaned:

```
from pyspark.sql.functions import regexp_replace, regexp_extract
from pyspark.sql.functions import col

regexp_string='years|year|\\+|\\<' #before + we nneed to prefix \\ or else it will take the + symbol as an operator.
df_sel.select(regexp_replace(col('emp_length'),regexp_string,"").alias("emplenght_cleaned"),col("emp_length")).show()
```

emplenght_cleaned	emp_length
1	< 1 year
10	10+ years
1	< 1 year
n/a	n/a
5	5 years
9	9 years
3	3 years
10	10+ years
n/a	n/a
10	10+ years
1	< 1 year
10	10+ years
10	10+ years
6	6 years
5	5 years
10	10+ years
2	2 years
10	10+ years
1	< 1 year
6	6 years

Cleaning the 'term' coumn from '36 months' to 36



```

+-----+-----+
|term_cleaned|      term|
+-----+-----+
|          36| 36 months|
|          36| 36 months|
|          60| 60 months|
|          60| 60 months|
|          60| 60 months|
|          36| 36 months|
|          36| 36 months|
|          36| 36 months|
|          36| 36 months|
|          36| 36 months|
|          60| 60 months|
|          60| 60 months|
|          60| 60 months|
|          36| 36 months|
|          60| 60 months|
|          60| 60 months|
|          36| 36 months|
|          60| 60 months|
|          60| 60 months|
|          36| 36 months|
+-----+-----+
only showing top 20 rows

```

✓ 0s completed at 12:58 PM

In the above one still it is not assigned to df. After assigning it to df,


```
df_sel.select("term","term_cleaned","emp_length","emplen_cleaned").show()
```

term	term_cleaned	emp_length	emplen_cleaned
36 months	36	< 1 year	1
36 months	36	10+ years	10
60 months	60	< 1 year	1
60 months	60	n/a	
60 months	60	5 years	5
36 months	36	9 years	9
36 months	36	3 years	3
36 months	36	10+ years	10
36 months	36	n/a	
36 months	36	10+ years	10
60 months	60	< 1 year	1
60 months	60	10+ years	10
60 months	60	10+ years	10
36 months	36	6 years	6
60 months	60	5 years	5
60 months	60	10+ years	10
36 months	36	2 years	2
60 months	60	10+ years	10
60 months	60	< 1 year	1
36 months	36	6 years	6

only showing top 20 rows

8. Print the schema



```
df_sel.printSchema()
```

```
root
```

```
 |-- term: string (nullable = true)
 |-- home_ownership: string (nullable = true)
 |-- grade: string (nullable = true)
 |-- purpose: string (nullable = true)
 |-- int_rate: string (nullable = true)
 |-- addr_state: string (nullable = true)
 |-- loan_status: string (nullable = true)
 |-- application_type: string (nullable = true)
 |-- loan_amnt: integer (nullable = true)
 |-- emp_length: string (nullable = true)
 |-- annual_inc: double (nullable = true)
 |-- dti: double (nullable = true)
 |-- delinq_2yrs: integer (nullable = true)
 |-- revol_util: string (nullable = true)
 |-- total_acc: integer (nullable = true)
 |-- num_tl_90g_dpd_24m: integer (nullable = true)
 |-- dti_joint: double (nullable = true)
 |-- term_cleaned: string (nullable = true)
 |-- emplen_cleaned: string (nullable = true)
```

Even though the the column is cleaned and it has only the numerical values spark takes this as '**string**' to convert the datatype into '**int**'we need to do it manually(Redefine the schema).

Correlation Matrix:

1. Checking for Correlation.

```
0s  '''df_sel.stat.corr('annual_inc','loan_amnt)''' #stat-refers to statistic function,corr-correlation between two column
#In here we are checking if the two columns are correlated to each other
#-1- Negatively correlated
#+1 - Positively correlated
#0 - No Correlation

#OR
spark.sql("select corr(annual_inc,loan_amnt).from loanstatus_sel").show() #in sql
```

```
+-----+
|corr(annual_inc, loan_amnt)|
+-----+
|          0.20103225337914624|
+-----+
```

[+ Code](#) [+ Text](#)

```
[ ] #from above one we can see it is not so much correlated.
```

```
1s [27] spark.sql("select corr(loan_amnt,term_cleaned) from loanstatus_sel").show()
```

```
+-----+
|corr(loan_amnt, term_cleaned)|
+-----+
|          0.3925941141844244|
+-----+
```

```
[ ] #still not a strong correlation
```

2. Cross Tab

```
7s  #Crosstab - displays the frequency distribution of values for
df_sel.stat.crosstab('loan_status','grade').show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
| loan_status_grade|    A|    B|    C|    D|    E|    F|    G|
+-----+-----+-----+-----+-----+-----+-----+
| Fully Paid| 1188| 1333| 1175|  807| 360|  36|   9|
| In Grace Period|  74|  112|  146|  122|  54|   4|   1|
| Charged Off|  16|   35|   38|   19|   8|   3|   3|
| Late (31-120 days)|  68|  164|  234|  220| 142|  14|   7|
| Current|36639|34139|29323|15823|5352|321|  79|
| Late (16-30 days)|  26|   78|  102|   81|  46|   9|   2|
+-----+-----+-----+-----+-----+-----+-----+
```

3. Frequency

```
1s  freq=df_sel.stat.freqItems(['purpose','grade'],0.3)
freq.collect()
```

```
[Row(purpose_freqItems=['debt_consolidation', 'credit_card', 'other'], grade_freqItems=['A', 'B', 'C'])]
```

Aggregate Functions:

1. Checking how many null values are present in the column

```
#Checking how many null values are present in the column
from pyspark.sql.functions import isnan, when, count, col
df_sel.select([count(when(isnan(c) | col(c).isNull(), c)).alias(c) for c in df_sel.columns]).show()
```

term	home_ownership	grade	purpose	int_rate	addr_state	loan_status	application_type	loan_amnt	emp_length	annual_inc	dti	delinq_2yrs	revol_util	total_acc	num_tl_90g_dpd_24m	dti_joint
0	0	0	0	0	0	0	0	0	0	0	237	0	156	0	0	111630

Cleaning the 'revol_util' column

```
#DTI Column
df_sel.describe("dti", "revol_util").show()
```

summary	dti	revol_util
count	128175	128256
mean	19.933177530719778	null
stddev	20.143542243475476	null
min	0.0	0%
max	999.0	99.90%

Analysis:

```
✓ 5s #DTI Column
df_sel.describe("dti","revol_util").show()
```

summary	dti	revol_util
count	128175	128256
mean	19.933177530719778	null
stddev	20.143542243475476	null
min	0.0	0%
max	999.0	99.90%

In the above image it shows the max value as 99 but in depth when analyzed,

```
✓ 4s from pyspark.sql.functions import regexp_replace, regexp_extract
#spark.sql("select ceil(regexp_replace(revol_util,'%','')), count(*) from loanstatus_sel group by ceil(regexp_replace(revol_util,'%',''))")
spark.sql("select ceil(regexp_replace(revol_util,'%','')), count(*) from loanstatus_sel group by ceil(regexp_replace(revol_util,'%',''))").show()
```

CEIL(regexp_replace(revol_util, '%', ''))	count(1)
29	1824
26	1776
65	1297
54	1582
19	1537
0	1132
112	2
22	1665
7	944
77	964
34	1909
184	1
126	1
94	536
50	1697
110	3
57	1586
32	1766
43	1729
84	794

only showing top 20 rows

[]

✓ 4s completed at 10:35 AM

It shows the max value as 184 but we can get only as 99. The reason behind this is it takes this as a **String[So that 9>1]**. Now we need to convert it into **int**.

After cleaning we got the below output

```
[26] df_sel=df_sel.withColumn("revolutil_cleaned",regexp_extract(col("revol_util"),"\\d+",0))

df_sel.describe('revol_util','revolutil_cleaned').show()
```

summary	revol_util	revolutil_cleaned
count	128256	128256
mean	null	43.76206961077844
stddev	null	24.801849528207015
min	0%	0
max	99.90%	99

'revol_util' is the column before cleaned. Once cleaned we can see the 'mean&stddev' values in it.

But still the max value is 99 since we haven't converted it into 'double'/'int'.
Filling the null vales with average values present in it.

```
#Defining a function since the columns may conatin a lot of null values in it
def fill_avg(df,colname):
    return df.select(colname).agg(avg(colname))

rev_avg=fill_avg(df_sel,'revolutil_cleaned')
```

Using Coalesce function.

```
#coalesce
...
It will take two variables
-if the first is not null it will take it.
-if the first is null it will take the second variable
...
from pyspark.sql.functions import coalesce
df_sel=df_sel.withColumn('revolutil_cleaned',coalesce(col('revolutil_cleaned'),col('rev_avg')))
...
In the above one['revolutil_cleaned'] it will take all the not null values and wherever there is null it will use the 'rev_avg'.
So where ever there is a null it will substitute with the average(will convert the two columns into single columns)
...
```

Converting the 'revolutil_cleaned' column to 'double' and displaying the describe function.

```

[38] #converting the datatype into double for the column 'revolutil_cleaned'
df_sel=df_sel.withColumn('revolutil_cleaned',df_sel['revolutil_cleaned'].cast('double'))

df_sel.describe('revol_util','revolutil_cleaned').show()

```

summary	revol_util	revolutil_cleaned
count	128256	128412
mean	null	43.762069610778525
stddev	null	24.786779696453955
min	0%	0.0
max	99.90%	183.0

From the above picture we could able to see now **max** is showing the correct value 183.0 where the old value is 99.90%.

DTI Column:

```

#checking how many null values are present in the column
from pyspark.sql.functions import isnan,when,count,col
df_sel.select((count(when(isnan(c)|col(c).isNull(),c)).alias(c) for c in df_sel.columns)).show()

```

term	home_ownership	grade	purpose	int_rate	addr_state	loan_status	application_type	loan_amnt	emp_length	annual_inc	dti	delinq_2yrs	revol_util	total_acc	num_tl_90p_dpd_24m	dti_joint
0	0	0	0	0	0	0	0	0	0	0	237	0	156	0	0	111630

Still DTI column is showing it has 237 null values, Lets check which values are null.

```

#checking which values are null in the 'dti' column
spark.sql("select * from loanstatus_sel where dti is null").show()

```

term	home_ownership	grade	purpose	int_rate	addr_state	loan_status	application_type	loan_amnt	emp_length	annual_inc	dti	delinq_2yrs	revol_util	total_acc	num_tl_90p_dpd_24m	dti_joint
60 months	MORTGAGE	B	major_purchase	10.72%	AZ	Current	Joint App	13000	n/a	0.0	null	0	26%	47		
60 months	RENT	C	debt_consolidation	16.91%	CA	Current	Joint App	18000	n/a	0.0	null	0	35.20%	12		
60 months	MORTGAGE	C	debt_consolidation	16.91%	NY	Fully Paid	Joint App	35000	n/a	0.0	null	0	90.10%	39		
36 months	RENT	C	other	13.56%	CA	Current	Joint App	5500	n/a	0.0	null	0	13%	17		
36 months	MORTGAGE	B	debt_consolidation	10.33%	TX	Current	Joint App	4700	n/a	0.0	null	0	4.40%	15		
36 months	RENT	A	debt_consolidation	7.56%	UT	Current	Joint App	10000	n/a	0.0	null	0	77.80%	20		
36 months	MORTGAGE	B	debt_consolidation	10.72%	CA	Current	Joint App	6000	10+ years	0.0	null	1	39%	22		
36 months	MORTGAGE	C	debt_consolidation	14.47%	GA	Current	Joint App	5000	n/a	0.0	null	1	86.20%	15		
36 months	RENT	E	debt_consolidation	23.40%	NY	Current	Joint App	40000	n/a	0.0	null	0	93.90%	27		
36 months	RENT	D	credit_card	18.94%	CA	Current	Joint App	35000	n/a	0.0	null	0	60.50%	15		
36 months	MORTGAGE	A	credit_card	7.56%	NJ	Current	Joint App	20000	n/a	0.0	null	0	65.20%	19		
36 months	MORTGAGE	B	debt_consolidation	11.31%	TN	Current	Joint App	30000	n/a	0.0	null	0	75.80%	19		
60 months	MORTGAGE	D	debt_consolidation	20.89%	AZ	Current	Joint App	28000	n/a	0.0	null	0	6%	30		
36 months	MORTGAGE	E	debt_consolidation	25.34%	KS	Current	Joint App	40000	n/a	0.0	null	0	97.50%	9		
60 months	MORTGAGE	D	debt_consolidation	17.97%	CO	Current	Joint App	10000	n/a	0.0	null	3	76%	26		
36 months	RENT	B	debt_consolidation	11.80%	CA	Current	Joint App	12000	n/a	0.0	null	0	100.40%	9		
36 months	MORTGAGE	A	other	7.56%	CA	Current	Joint App	7000	n/a	0.0	null	0	3.70%	16		
60 months	MORTGAGE	C	debt_consolidation	15.02%	FL	Current	Joint App	10000	n/a	0.0	null	0	98.30%	10		
60 months	MORTGAGE	B	credit_card	10.72%	SC	Current	Joint App	36000	n/a	0.0	null	0	13.70%	9		
60 months	MORTGAGE	D	debt_consolidation	18.94%	NY	Current	Joint App	16000	n/a	0.0	null	0	49.50%	54		

When check we can see wherever the **application_type='Joint App'** the dti value is null.

Looking for Insights:

```
spark.sql("select application_type,dti,dti_joint from loanstatus_sel where dti is null").show()
```

application_type	dti	dti_joint
Joint App	null	33.06
Joint App	null	17.67
Joint App	null	27.3
Joint App	null	8.74
Joint App	null	11.68
Joint App	null	28.01
Joint App	null	15.06
Joint App	null	12.79
Joint App	null	8.05
Joint App	null	24.22
Joint App	null	20.04
Joint App	null	12.75
Joint App	null	24.89
Joint App	null	22.61
Joint App	null	19.01
Joint App	null	15.34
Joint App	null	0.6
Joint App	null	23.69
Joint App	null	17.29
Joint App	null	29.25

only showing top 20 rows

From the above insights, we are going to create a new column name 'dti_cleaned' in that with the help of coalesce function, 'dti' records null values will be filled with the values of 'dti_joint'.

```
#from the above insights ,
...
we are going to create a new column name 'dti_cleaned' in that with the help of coalesce function, 'dti' records null values will be filled with the vlaues of 'dti_joint'
...
df_sel=df_sel.withColumn("dti_cleaned",coalesce(col("dti"),col("dti_joint")))
#checking the result
from pyspark.sql.functions import isnan,when,count,col
df_sel.select([count(when(isnan(c)|col(c).isNull(),c)).alias(c) for c in df_sel.columns]).show()
```

nt_rate	addr_state	loan_status	application_type	loan_amnt	emp_length	annual_inc	dti	delinq_2yrs	revol_util	total_acc	num_tl_90g_dpd_24m	dti_joint	revolutil_cleaned	rev_avg	dti_cleaned
0	0	0	0	0	0	0	237	0	156	0	111630	0	0	0	0

Cleaning 'loan_status'

```
df_sel.groupby('loan_status').count().show()
```

loan_status	count
Fully Paid	4908
In Grace Period	513
Charged Off	122
Late (31-120 days)	849
Current	121676
Late (16-30 days)	344

From this we can understand only 'Fully Paid' alone shows that it has full paid, other values are showing like the payment is delayed then **Bad loan**.

```
#adding a new column as bad loan
from pyspark.sql import SparkSession
from pyspark.sql.functions import when
df_sel=df_sel.withColumn("bad loan",when(df_sel.loan_status.isin(["Late (31-120 days)","Charged Off","In Grace Period","Late (16-30 days)"]),'Yes').otherwise("No"))

df_sel.groupby('bad loan').count().show()
```

bad loan	count
No	126584
Yes	1828

From this we can see majority of loan is good loan.

When checking on the bad loan we can see the intrest rate are pretty high.

```
df_sel.filter(df_sel.bad_loan == 'Yes').show()
```

term	home_ownership	grade	purpose	int_rate	addr_state	loan_status	application_type	loan_amnt	emp_length	annual_inc	dti	delinq_2yrs	revol_util	total_acc	num_t
36 months	RENT	B	debt consolidation	10.33%	PA	Late (31-120 days)	Individual	16000	3 years	71000.0	16.58	0	11.80%	28	
60 months	MORTGAGE	D	debt consolidation	22.35%	OH	Late (31-120 days)	Joint App	17500	9 years	50000.0	36.92	0	35.40%	37	
60 months	RENT	C	debt consolidation	15.02%	NY	Late (31-120 days)	Individual	30000	2 years	90000.0	14.28	0	14.90%	10	
36 months	MORTGAGE	A	debt consolidation	8.19%	PA	Late (31-120 days)	Individual	20975	2 years	165000.0	24.27	0	16.60%	21	
60 months	RENT	A	debt consolidation	8.19%	CA	Late (31-120 days)	Individual	14000	< 1 year	65000.0	13.48	0	39.10%	18	
36 months	RENT	A	debt consolidation	8.19%	WA	Late (31-120 days)	Individual	14000	1 year	36000.0	25.83	0	73.60%	11	
36 months	RENT	B	vacation	11.80%	NC	Late (31-120 days)	Individual	20000	3 years	50000.0	9.18	0	13.90%	26	
36 months	RENT	B	debt consolidation	11.80%	NJ	Late (31-120 days)	Individual	32000	< 1 year	50000.0	6.1	0	50.30%	28	
36 months	OWN	E	other	26.31%	IN	Late (31-120 days)	Individual	9100	3 years	62000.0	13.38	0	12.10%	23	
36 months	RENT	C	credit card	13.56%	OR	Late (31-120 days)	Individual	5000	1 year	50000.0	16.77	0	30.50%	9	
60 months	RENT	C	major purchase	14.47%	CA	Late (31-120 days)	Individual	36000	9 years	93600.0	4.14	0	0.30%	29	
36 months	RENT	E	credit card	24.37%	NY	Late (31-120 days)	Individual	19500	10+ years	56000.0	29.83	0	66%	33	
36 months	RENT	E	vacation	25.34%	NY	Late (31-120 days)	Individual	1500	< 1 year	20000.0	9.24	0	83.20%	3	
36 months	MORTGAGE	E	debt consolidation	25.34%	GA	Late (31-120 days)	Individual	3525	10+ years	150000.0	14.25	0	31.10%	15	
60 months	MORTGAGE	C	debt consolidation	16.91%	CA	Late (31-120 days)	Individual	11000	10+ years	97000.0	28.79	0	37.70%	32	
36 months	RENT	C	debt consolidation	13.56%	MT	Late (31-120 days)	Individual	11500	8 years	33000.0	2.55	0	20%	12	
36 months	MORTGAGE	A	debt consolidation	8.81%	CA	Late (31-120 days)	Individual	22000	5 years	165000.0	15.59	0	46.40%	46	
60 months	OWN	D	debt consolidation	17.97%	NV	Late (31-120 days)	Individual	19200	9 years	85000.0	39.11	4	28.30%	53	
60 months	RENT	B	debt consolidation	11.80%	IN	Late (31-120 days)	Individual	40000	5 years	130000.0	17.88	0	78.30%	16	
36 months	MORTGAGE	A	credit card	6.46%	WA	Charged Off	Individual	8000	10+ years	55000.0	17.83	0	29.10%	49	

only showing top 20 rows

Dropping the column:

```
✓ [19] #Dropping the unnecessary columns
0s df_sel_final=df_sel.drop('revol_util','dti','dti_joint','bad_loan')

✓ df_sel_final.printSchema()
ps
root
|-- term: string (nullable = true)
|-- home_ownership: string (nullable = true)
|-- grade: string (nullable = true)
|-- purpose: string (nullable = true)
|-- int_rate: string (nullable = true)
|-- addr_state: string (nullable = true)
|-- loan_status: string (nullable = true)
|-- application_type: string (nullable = true)
|-- loan_amnt: integer (nullable = true)
|-- emp_length: string (nullable = true)
|-- annual_inc: double (nullable = true)
|-- delinq_2yrs: integer (nullable = true)
|-- total_acc: integer (nullable = true)
|-- num_tl_90g_dpd_24m: integer (nullable = true)
|-- bad_loan: string (nullable = false)
```

Storing data in a permanent table

```
▶ #permanent table
#saving table in the parquet format
permanent_table='lc_loan_data'
df_sel.write.format("parquet").saveAsTable(permanent_table)
```

Notes:

1. Data frames in spark works in a distributed format like they will divide into chunks of parts and they assign it to worker node.
2. Data frames in Pandas is mostly dependent on Memory.
3. Data frames in spark has more function and used well for analyzing than SQL.
4. Spark is an lazy evaluation framework creates a DAG and then an action command comes and executes the DAG. So that it optimizes the DAG in the distributed way.
5. Cache()-Stores the df in the memory , subsequent operation can be performed faster without having to recompute the dataframe from source. If we cache too many dataframes will lead to excessive memory consumption and cause memory related issue.